

Predictive Modeling Report

Giodanno Limin

2025-11-27

Contents

```
# get the training data
train <- read.csv("trainingdata.csv")

# direct access
attach(train)
# write a cv_rmse function
cv_rmse <- function(formula, data, K = 10) {
  set.seed(1)
  n <- nrow(data)
  folds <- sample(rep(1:K, length.out = n))

  rmse <- numeric(K)

  for (k in 1:K) {
    train_fold <- data[folds != k, ]
    test_fold <- data[folds == k, ]

    fit <- lm(formula, data=train_fold)
    pred <- predict(fit, newdata=test_fold)

    rmse[k] <- sqrt(mean((pred - test_fold$y)^2))
  }

  return(list(
    mean_RMSE = mean(rmse),
    sd_RMSE   = sd(rmse),
    all_RMSE   = rmse
  ))
}

lm1 = lm(y~.,data=train)
summary(lm1)

## Call:
## lm(formula = y ~ ., data = train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.40482 -0.36592 -0.04287  0.27987  2.08131
```

```

##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -1.836e+00 1.285e+00 -1.428 0.154835
## X1          -1.244e-01 1.741e-01 -0.714 0.475945
## X2          -3.155e-02 1.675e-01 -0.188 0.850772
## X3          -5.407e-03 1.359e-02 -0.398 0.691125
## X4          1.087e-01 1.669e-01  0.651 0.515867
## X5          5.496e-02 1.750e-01  0.314 0.753820
## X6          -1.084e-02 1.814e-01 -0.060 0.952411
## X7          -1.069e-01 4.619e-02 -2.314 0.021754 *
## X8          1.281e-02 4.768e-02  0.269 0.788477
## X9          -9.935e-02 4.636e-02 -2.143 0.033388 *
## X10         1.234e-02 4.846e-02  0.255 0.799310
## X11         7.327e-03 4.632e-02  0.158 0.874471
## X12         2.535e-01 1.700e-01  1.491 0.137711
## X13         -2.748e-01 1.676e-01 -1.640 0.102688
## X14         1.093e-01 1.584e-01  0.690 0.491207
## X15         6.753e-02 1.627e-01  0.415 0.678518
## X16         1.821e-01 1.702e-01  1.070 0.285863
## X17         8.294e-02 4.769e-02  1.739 0.083638 .
## X18         -2.416e-02 5.150e-02 -0.469 0.639496
## X19         8.760e-02 4.314e-02  2.030 0.043730 *
## X20         -3.196e-02 1.651e-01 -0.194 0.846704
## X21         4.557e-02 7.392e-03  6.165 4.24e-09 ***
## X22         -9.078e-03 4.641e-02 -0.196 0.845136
## X23         1.773e-02 4.684e-02  0.379 0.705401
## X24         -2.134e-02 4.716e-02 -0.452 0.651525
## X25         -1.274e-01 1.719e-01 -0.741 0.459598
## X26         1.752e-02 1.637e-01  0.107 0.914886
## X27         -1.039e-01 1.659e-01 -0.627 0.531665
## X28         1.733e-01 1.642e-01  1.055 0.292700
## X29         1.317e-02 4.700e-02  0.280 0.779661
## X30         -9.838e-02 3.017e-02 -3.260 0.001323 **
## X31         9.855e-02 1.668e-01  0.591 0.555418
## X32         1.385e-02 1.072e-01  0.129 0.897335
## X33         -2.752e-02 1.629e-01 -0.169 0.866020
## X34         1.455e-01 1.681e-01  0.866 0.387872
## X35         -2.209e-01 1.736e-01 -1.272 0.204831
## X36         -2.453e-03 3.882e-03 -0.632 0.528307
## X37         -1.959e-02 1.665e-01 -0.118 0.906462
## X38         1.836e-05 1.013e-04  0.181 0.856392
## X39         -1.197e-01 1.564e-01 -0.765 0.445066
## X40         4.388e-01 1.685e-01  2.605 0.009937 **
## X41         1.551e-01 1.147e-01  1.351 0.178233
## X42         1.318e-02 4.826e-02  0.273 0.785022
## X43         1.679e-01 1.708e-01  0.983 0.326747
## X44         3.610e-02 1.737e-01  0.208 0.835568
## X45         -3.232e-02 7.625e-02 -0.424 0.672183
## X46         -2.032e-01 1.713e-01 -1.187 0.236865
## X47         3.214e-02 1.660e-01  0.194 0.846703
## X48         -1.260e-02 2.972e-02 -0.424 0.672116
## X49         -5.978e-02 1.653e-01 -0.362 0.718031
## X50         1.771e-02 1.614e-01  0.110 0.912759

```

## X51	8.932e-02	1.806e-01	0.494	0.621544
## X52	-1.835e-01	1.766e-01	-1.039	0.300114
## X53	3.609e-02	1.730e-01	0.209	0.835017
## X54	1.805e-02	4.775e-02	0.378	0.705803
## X55	-1.776e-01	1.795e-01	-0.989	0.323782
## X56	1.639e-01	1.684e-01	0.973	0.331777
## X57	-9.789e-02	1.711e-01	-0.572	0.568028
## X58	2.930e-02	4.481e-02	0.654	0.514000
## X59	-7.999e-02	6.554e-02	-1.220	0.223814
## X60	-6.243e-02	4.555e-02	-1.370	0.172196
## X61	-4.652e-02	4.842e-02	-0.961	0.337910
## X62	-1.158e-02	1.774e-01	-0.065	0.948023
## X63	9.788e-02	1.741e-01	0.562	0.574619
## X64	-6.892e-02	4.710e-02	-1.463	0.145115
## X65	1.156e-02	2.238e-01	0.052	0.958841
## X66	5.121e-02	4.589e-02	1.116	0.265862
## X67	2.353e-01	1.613e-01	1.459	0.146250
## X68	1.476e-01	1.658e-01	0.891	0.374203
## X69	-1.332e-02	4.877e-02	-0.273	0.785088
## X70	4.670e-02	4.646e-02	1.005	0.316192
## X71	7.835e-02	4.736e-02	1.654	0.099722 .
## X72	-4.705e-02	1.755e-01	-0.268	0.788985
## X73	5.781e-02	4.611e-02	1.254	0.211556
## X74	-1.121e-01	1.801e-01	-0.622	0.534543
## X75	-4.404e-02	1.714e-01	-0.257	0.797470
## X76	-6.731e-03	1.619e-01	-0.042	0.966891
## X77	-4.527e-03	4.596e-02	-0.099	0.921640
## X78	-3.080e-01	1.635e-01	-1.884	0.061144 .
## X79	-5.365e-02	1.701e-01	-0.315	0.752839
## X80	5.361e-02	4.567e-02	1.174	0.241975
## X81	-1.099e-02	4.608e-02	-0.238	0.811805
## X82	-1.016e-01	1.603e-01	-0.634	0.527028
## X83	3.613e-02	4.447e-02	0.812	0.417566
## X84	-5.795e-02	1.678e-01	-0.345	0.730214
## X85	7.753e-02	1.779e-01	0.436	0.663545
## X86	-3.172e-01	1.649e-01	-1.924	0.055922 .
## X87	-3.311e-01	1.608e-01	-2.059	0.040895 *
## X88	1.053e-01	4.483e-02	2.350	0.019817 *
## X89	5.708e-02	4.830e-02	1.182	0.238821
## X90	-2.541e-02	4.021e-02	-0.632	0.528302
## X91	8.141e-03	4.964e-02	0.164	0.869909
## X92	-1.035e-02	1.637e-01	-0.063	0.949637
## X93	1.041e-02	4.654e-02	0.224	0.823168
## X94	-1.596e-02	1.785e-01	-0.089	0.928847
## X95	2.680e-01	1.637e-01	1.637	0.103335
## X96	2.732e-03	4.564e-02	0.060	0.952322
## X97	2.712e-01	1.783e-01	1.521	0.129900
## X98	2.004e-01	1.718e-01	1.166	0.245107
## X99	-1.650e-01	1.674e-01	-0.986	0.325498
## X100	4.192e-02	4.946e-02	0.848	0.397738
## X101	2.041e-01	1.671e-01	1.222	0.223433
## X102	-9.294e-03	5.394e-02	-0.172	0.863376
## X103	2.856e-02	4.616e-02	0.619	0.536885
## X104	1.412e-02	1.551e-01	0.091	0.927562

```

## X105      -6.734e-03  1.711e-01 -0.039  0.968653
## X106     -1.323e-02  4.811e-02 -0.275  0.783576
## X107     -1.684e-02  4.641e-02 -0.363  0.717120
## X108    -9.836e-05  7.310e-05 -1.346  0.180074
## X109    -1.411e-01  1.734e-01 -0.814  0.416648
## X110    -1.984e-01  1.583e-01 -1.253  0.211699
## X111     1.065e+00  2.877e-01  3.702  0.000281 ***
## X112    -4.978e-04  4.910e-02 -0.010  0.991921
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.6459 on 187 degrees of freedom
## Multiple R-squared:  0.6955, Adjusted R-squared:  0.5131
## F-statistic: 3.814 on 112 and 187 DF,  p-value: 3.227e-16
cv_rmse(lm1,train)

## $mean_RMSE
## [1] 0.829424
##
## $sd_RMSE
## [1] 0.1684319
##
## $all_RMSE
## [1] 0.9346653 0.9864522 0.9430970 0.8060121 0.7775150 0.3984574 0.7934964
## [8] 0.9436986 0.8869845 0.8238614

sig_vars <- summary(lm1)$coefficients
sig_vars[sig_vars[, "Pr(>|t|)"] < 0.05, ]

##           Estimate Std. Error t value Pr(>|t|)
## X7     -0.10687732 0.046186860 -2.314020 2.175352e-02
## X9     -0.09935425 0.046357655 -2.143211 3.338833e-02
## X19     0.08759752 0.043142547  2.030421 4.373009e-02
## X21     0.04557020 0.007391737  6.165019 4.243496e-09
## X30     -0.09837712 0.030174855 -3.260235 1.322716e-03
## X40     0.43883711 0.168482914  2.604639 9.936539e-03
## X87     -0.33108819 0.160812368 -2.058848 4.089515e-02
## X88     0.10534484 0.044828094  2.349974 1.981740e-02
## X111    1.06523359 0.287735588  3.702127 2.812017e-04

lm2 = lm(y~X7+X9+X19+X21+X30+X40+X87+X88+X111,data=train)
summary(lm2)

##
## Call:
## lm(formula = y ~ X7 + X9 + X19 + X21 + X30 + X40 + X87 + X88 +
##     X111, data = train)
##
## Residuals:
##       Min     1Q Median     3Q    Max 
## -2.2782 -0.3967 -0.1091  0.3479  2.3933 
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)    
## (Intercept) -2.524555   0.277289 -9.104   < 2e-16 ***

```

```

## X7      -0.056952  0.043327 -1.314   0.1897
## X9      -0.083958  0.041889 -2.004   0.0460 *
## X19     0.266808  0.027271  9.784   < 2e-16 ***
## X21     0.050279  0.006094  8.250   5.61e-15 ***
## X30     -0.053695  0.026520 -2.025   0.0438 *
## X40      0.279613  0.110431  2.532   0.0119 *
## X87     -0.271269  0.108995 -2.489   0.0134 *
## X88     0.048036  0.042130  1.140   0.2551
## X111    1.455985  0.264727  5.500   8.34e-08 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.7167 on 290 degrees of freedom
## Multiple R-squared:  0.4186, Adjusted R-squared:  0.4005
## F-statistic: 23.2 on 9 and 290 DF, p-value: < 2.2e-16
cv_rmse(lm2,train)

```

```

## $mean_RMSE
## [1] 0.72777796
##
## $sd_RMSE
## [1] 0.1247808
##
## $all_RMSE
## [1] 0.6820813 0.7446164 0.8649398 0.8521659 0.8207434 0.4932418 0.5423195
## [8] 0.7996716 0.7368539 0.7411627

```

Let's try better way

```

cor_with_y <- function(x) {
  cor(train$y, x)
}
cors <- sapply(train[, -1], cor_with_y)
head(sort(abs(cors), decreasing=TRUE), 10)

```

```

##      X66      X69      X21      X19      X45      X111      X41      X38
## 0.5700591 0.4227833 0.4184403 0.4107010 0.2733106 0.1926304 0.1697941 0.1653557
##      X3      X59
## 0.1652030 0.1417632

```

```

top10 <- names(sort(abs(cors), decreasing = TRUE))[1:10]
top10

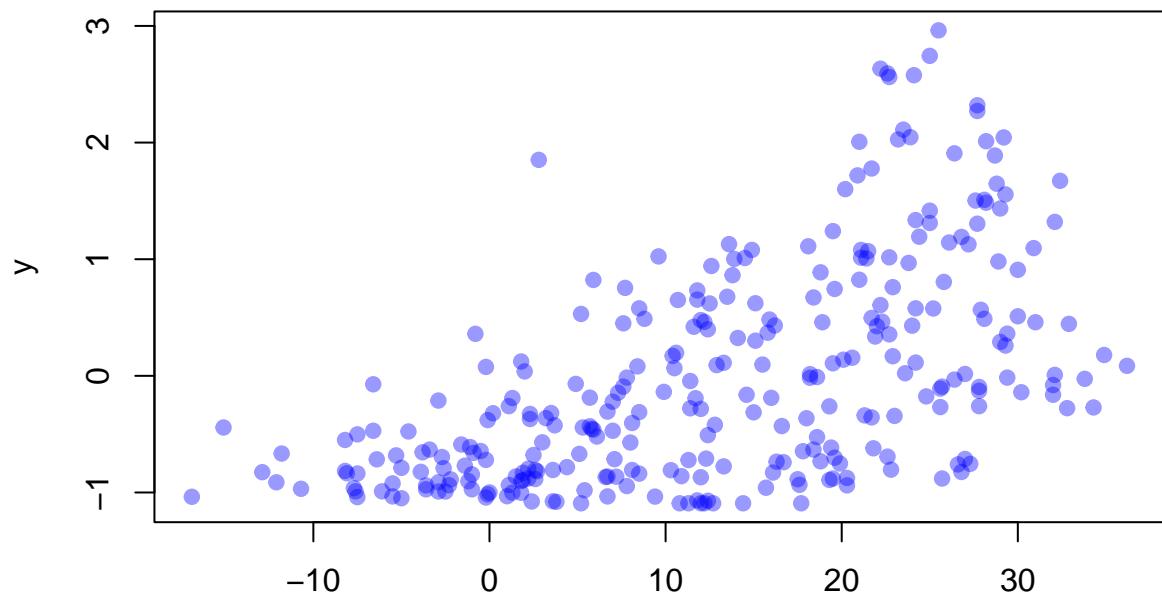
```

```

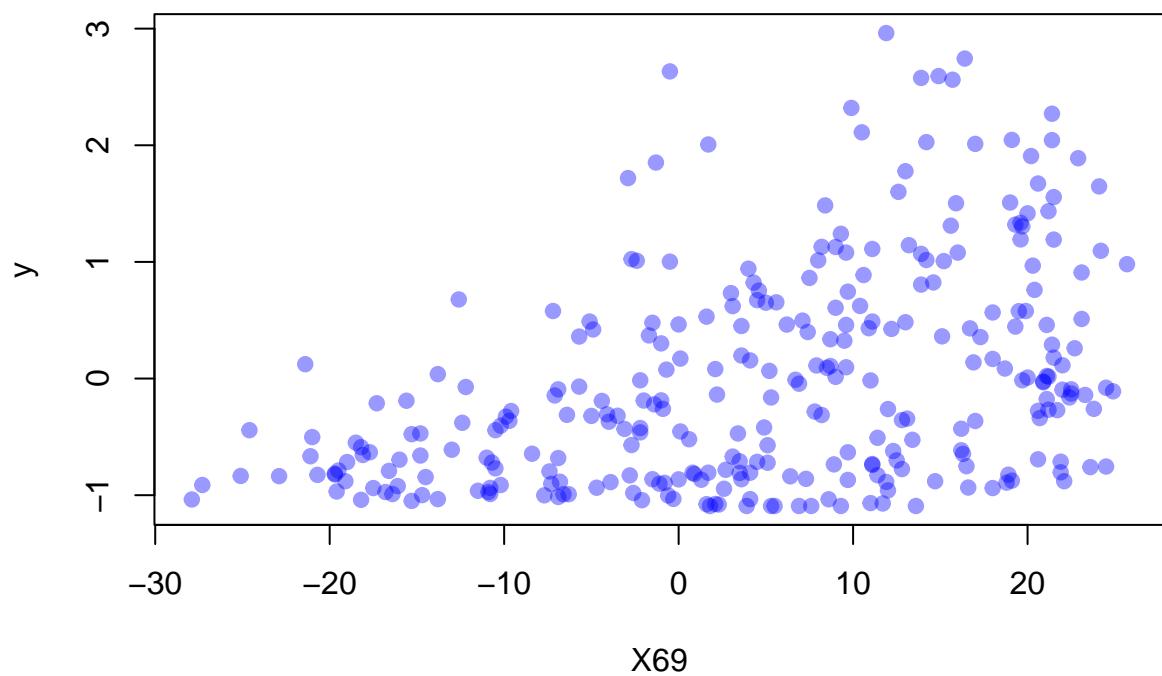
## [1] "X66"  "X69"  "X21"  "X19"  "X45"  "X111" "X41"  "X38"  "X3"   "X59"
for (var in top10) {
  plot(
    train[[var]], train$y,
    main = paste(var, "vs y"),
    xlab = var,
    ylab = "y",
    pch = 19,
    col = rgb(0, 0, 1, 0.4)
  )
}

```

X66 vs y

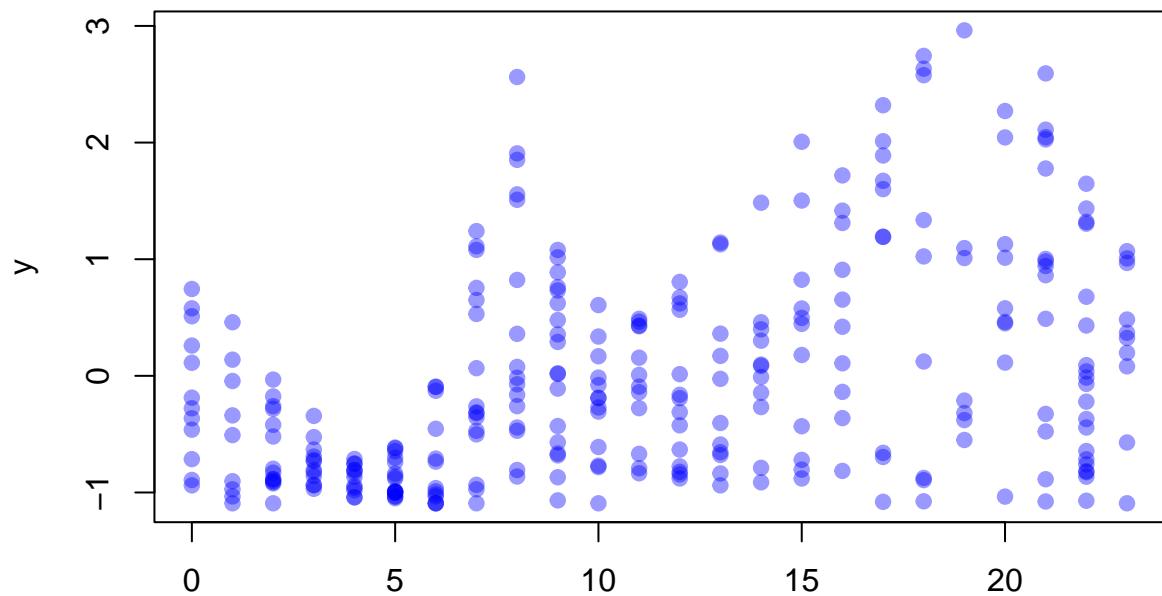


**X66
X69 vs y**

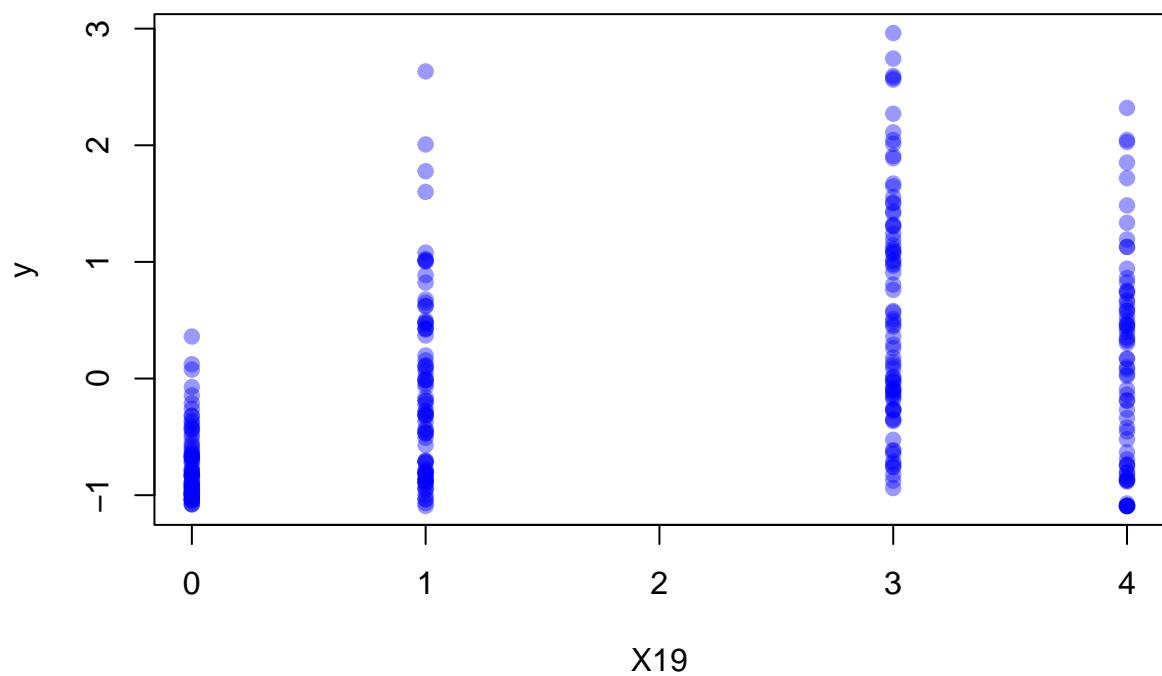


X69

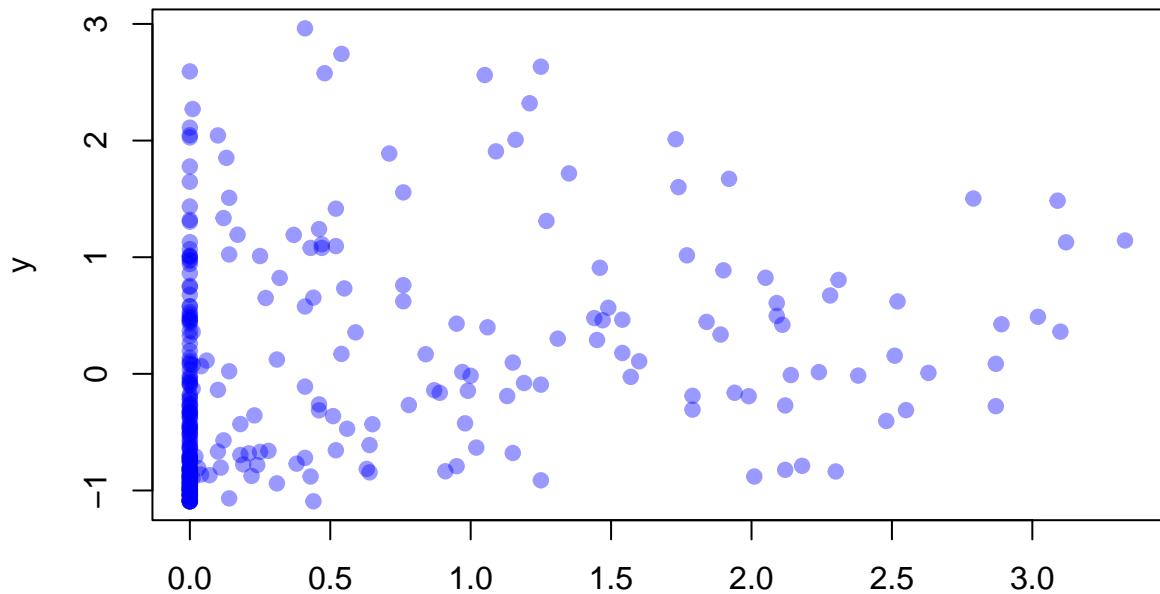
X21 vs y



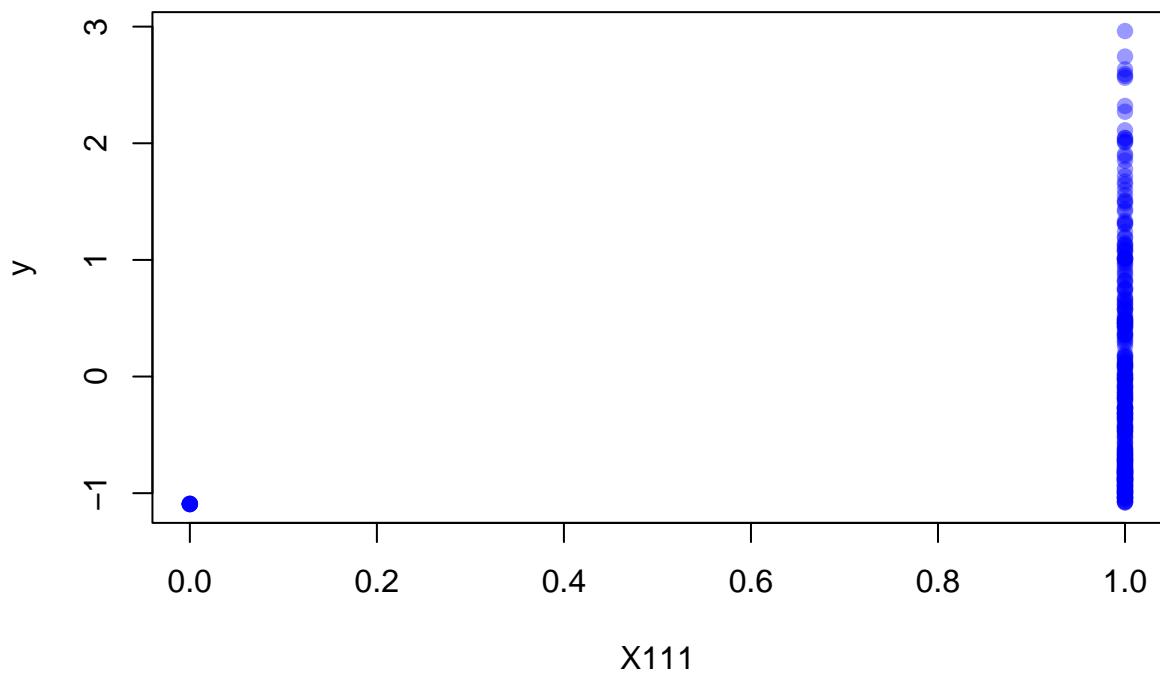
**X21
X19 vs y**



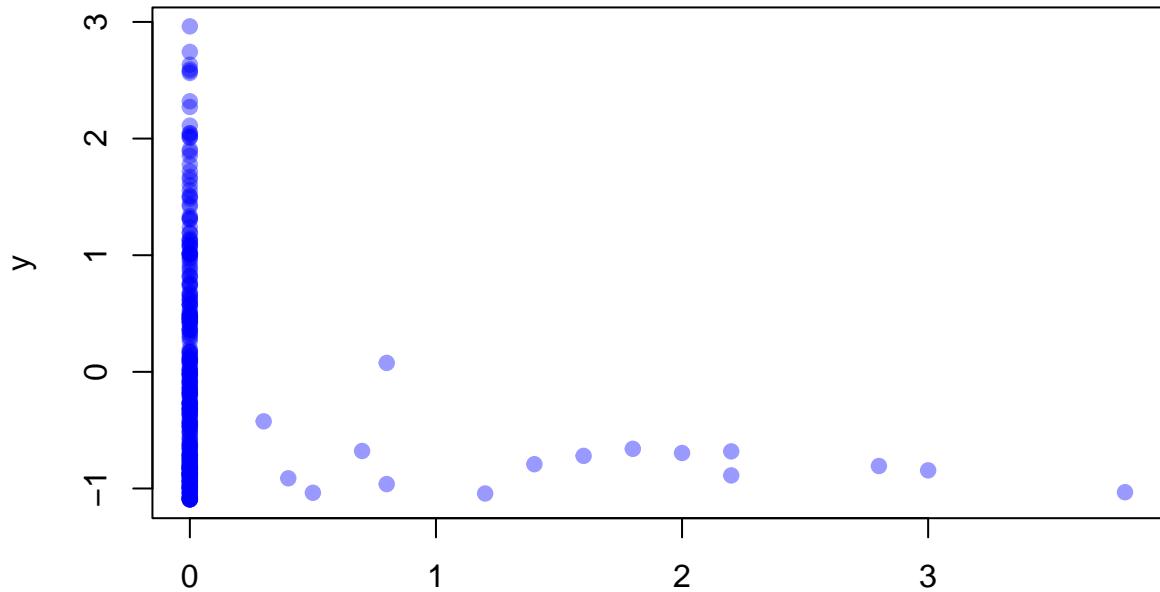
X45 vs y



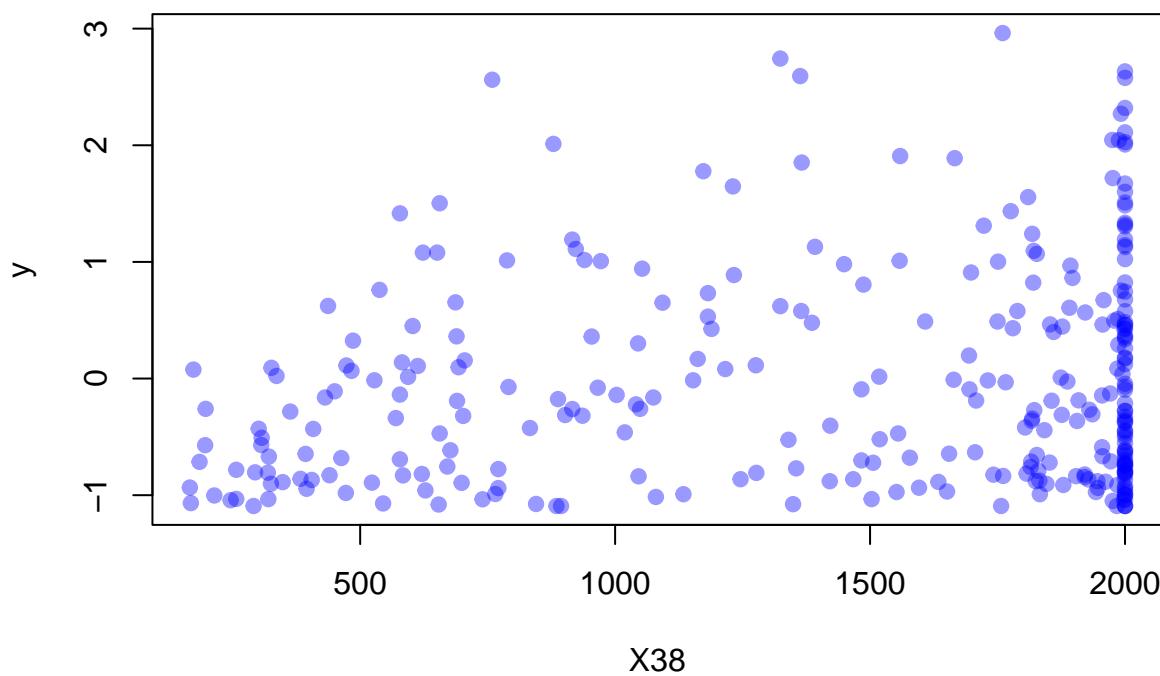
**X45
X111 vs y**



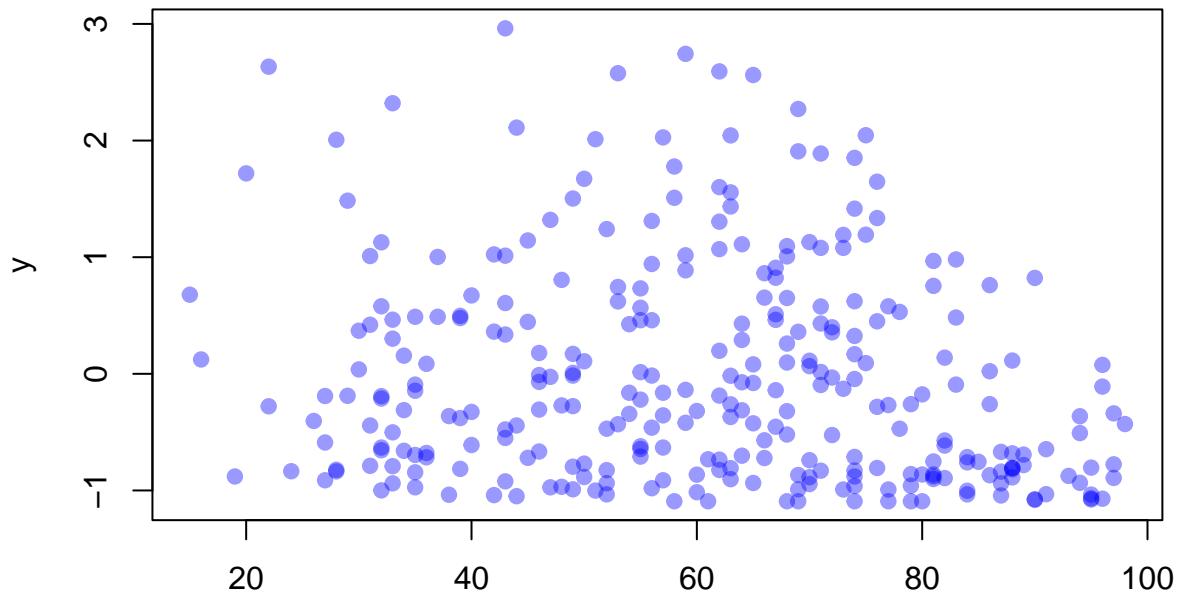
X41 vs y



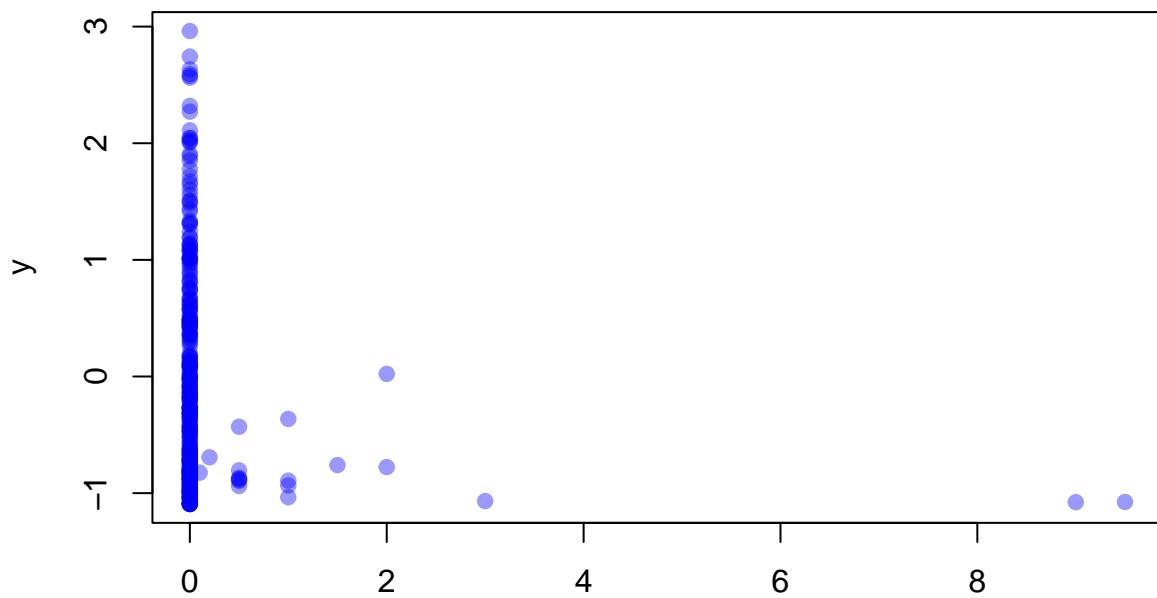
X41
X38 vs y



X3 vs y



X3
X59 vs y



X59

```
lm4 = lm(y~X66+X69+X21,data=train)
summary(lm4)
```

```
##
## Call:
## lm(formula = y ~ X66 + X69 + X21, data = train)
```

```

## 
## Residuals:
##      Min       1Q   Median      3Q      Max
## -1.58451 -0.42403 -0.09015  0.35334  2.56191
## 
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)    
## (Intercept) -1.294099  0.096381 -13.427 < 2e-16 ***
## X66          0.075799  0.008765   8.648 3.37e-16 ***
## X69          -0.033818  0.007909  -4.276 2.57e-05 ***
## X21          0.040972  0.005901   6.944 2.42e-11 ***
## ---    
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
## 
## Residual standard error: 0.6722 on 296 degrees of freedom
## Multiple R-squared:  0.4779, Adjusted R-squared:  0.4726 
## F-statistic: 90.33 on 3 and 296 DF,  p-value: < 2.2e-16

cv_rmse(lm4,train)

```

```

## $mean_RMSE
## [1] 0.6701781
## 
## $sd_RMSE
## [1] 0.09868585
## 
## $all_RMSE
## [1] 0.6039787 0.7512891 0.7767443 0.6238047 0.7737403 0.4714130 0.5797160
## [8] 0.7229067 0.6865790 0.7116088

```

Now let's check if the adding other variable lower the mean_RMSE "X19" "X45" "X111" "X41" "X38" "X3" "X59"

```

vars_to_test <- c("X19", "X45", "X111", "X41", "X38", "X3", "X59")

results <- data.frame(
  variable = vars_to_test,
  mean_RMSE = NA
)

base_form <- "y ~ X66 + X69 + X21"
base_rmse <- cv_rmse(as.formula(base_form), train)$mean_RMSE

results <- data.frame(
  variable = vars_to_test,
  mean_RMSE = NA,
  improvement = NA
)

for (i in seq_along(vars_to_test)) {

  var <- vars_to_test[i]

  # construct formula dynamically
  form <- as.formula(
    paste(base_form, "+", var)
}

```

```

}

# compute rmse
new_rmse <- cv_rmse(form, train)$mean_RMSE

# fill results
results$mean_RMSE[i] <- new_rmse
results$improvement[i] <- new_rmse < base_rmse
}

results

##   variable mean_RMSE improvement
## 1      X19  0.6672321      TRUE
## 2      X45  0.6709039     FALSE
## 3     X111  0.6555636      TRUE
## 4      X41  0.6724600     FALSE
## 5      X38  0.6712700     FALSE
## 6      X3   0.6672124      TRUE
## 7      X59  0.6641478      TRUE

lm5 = lm(y~X66+X69+X21+X19+X111+X3+X59,data=train)
cv_rmse(lm5,train)

## $mean_RMSE
## [1] 0.6411805
##
## $sd_RMSE
## [1] 0.1118216
##
## $all_RMSE
## [1] 0.6100652 0.6954469 0.7567000 0.6371614 0.7518331 0.4045187 0.4985694
## [8] 0.7063156 0.6701002 0.6810947

```

Now let's try polynomial for X66, X69, X21, and X3.

```

search_poly_multivar <- function(max66 = 10, max69 = 10, max21 = 10) {
  results <- data.frame()

  for (d66 in 1:max66) {
    for (d69 in 1:max69) {
      for (d21 in 1:max21) {

        form <- as.formula(
          paste(
            "y ~ poly(X66, ", d66, ", raw=TRUE) +",
            "poly(X69, ", d69, ", raw=TRUE) +",
            "poly(X21, ", d21, ", raw=TRUE) +",
            "+ X19 + X111 + X3 + X59",
            sep=" "
          )
    }

    cv <- cv_rmse(form, train)$mean_RMSE

    results <- rbind(

```

```

        results,
        data.frame(
          d66 = d66,
          d69 = d69,
          d21 = d21,
          RMSE = cv
        )
      )
    }
  }

  return(results)
}

# run the search
poly_results <- search_poly_multivar()
best_row <- poly_results[which.min(poly_results$RMSE), ]
best_row

##      d66 d69 d21      RMSE
## 430     5   3   10 0.5068829

search_poly_X3 <- function(max3=10) {
  results <- data.frame()

  for (d3 in 1:max3){

    form <- as.formula(
      paste(
        "y ~",
        "poly(X66, 5, raw=TRUE) +",
        "poly(X69, 3, raw=TRUE) +",
        "poly(X21, 10, raw=TRUE) +",
        "poly(X3, ", d3, ", raw=TRUE) +",
        "X19 + X111 + X59",
        sep=" "
      )
    )

    cv <- cv_rmse(form, train)$mean_RMSE

    results <- rbind(
      results,
      data.frame(
        d3 = d3,
        RMSE = cv
      )
    )
  }

  return(results)
}

```

```

# run the search
poly_results <- search_poly_X3()

best_row <- poly_results[ which.min(poly_results$RMSE) , ]
best_row

##   d3      RMSE
## 3  3 0.4961006

# let's check again for X21, since previously our max degree is 10,
# maybe if the degree greater than 10 is better

search_poly_X21 <- function(max21 = 20) {
  results <- data.frame()

  for (d21 in 10:max21) {

    form <- as.formula(
      paste(
        "y ~",
        "poly(X66, 5, raw=TRUE) +",
        "poly(X69, 3, raw=TRUE) +",
        "poly(X3, 3, raw=TRUE) +",
        "poly(X21, ", d21, ", raw=TRUE) +",
        "X19 + X111 + X59",
        sep=" "
      )
    )

    cv <- cv_rmse(form, train)$mean_RMSE

    results <- rbind(results, data.frame(d21=d21, RMSE=cv))
  }

  return(results)
}

poly21_results <- search_poly_X21(max21 = 15)

# best degree
best_row <- poly21_results[which.min(poly21_results$RMSE), ]
best_row

##   d21      RMSE
## 5  14 0.4931188

lm6 = lm(y~poly(X66, 5, raw=TRUE)+
          poly(X69, 3, raw=TRUE)+
          poly(X21, 14, raw=TRUE)+
          poly(X3, 3, raw=TRUE)+
          + X19 + X111 + X59)
cv_rmse(lm6,train)

## $mean_RMSE
## [1] 0.4931188
##

```

```

## $sd_RMSE
## [1] 0.07598946
##
## $all_RMSE
## [1] 0.5195799 0.5533502 0.5138166 0.4268541 0.5644166 0.4426720 0.3197027
## [8] 0.5090485 0.5285066 0.5532411

Now let's check the interactions for X66, X69, X21, X3

base_rmse <- cv_rmse(lm6, train)$mean_RMSE

base_string <- paste(
  "y ~",
  "poly(X66, 5, raw=TRUE) +",
  "poly(X69, 3, raw=TRUE) +",
  "poly(X21,14, raw=TRUE) +",
  "poly(X3, 3, raw=TRUE) +",
  "+ X19 + X111 + X59"
)

candidates <- c(
  "X66:X69",
  "X66:X21",
  "X66:X3",
  "X69:X21",
  "X69:X3",
  "X21:X3"
)

test_interactions <- function(base_string, candidates, train) {

  results <- data.frame()

  for (int in candidates) {
    full_form_string <- paste(base_string, "+", int)
    form <- as.formula(full_form_string)

    cv <- cv_rmse(form, train)$mean_RMSE

    results <- rbind(
      results,
      data.frame(
        interaction = int,
        RMSE = cv
      )
    )
  }

  return(results)
}

interaction_results <- test_interactions(base_string, candidates, train)
improved <- interaction_results$RMSE < base_rmse
interaction_results[ improved , ]

##   interaction     RMSE

```

```

## 2      X66:X21 0.4737751
## 4      X69:X21 0.4850535
## 6      X21:X3 0.4928832

lm7 = lm(y~poly(X66, 5, raw=TRUE) +
          poly(X69, 3, raw=TRUE) +
          poly(X21, 14, raw=TRUE) +
          poly(X3, 3, raw=TRUE) +
          +X66:X21+X69:X21+X21:X3+
          +X19+X111+X59)
cv_rmse(lm7,train)

## $mean_RMSE
## [1] 0.461731
##
## $sd_RMSE
## [1] 0.06842499
##
## $all_RMSE
## [1] 0.4493378 0.5376595 0.5001442 0.4012809 0.4938008 0.4157965 0.3281635
## [8] 0.4936212 0.4423954 0.5551105

```

Next let's try 2nd degree of interactions ()

```

generate_interactions2 <- function(var1, var2) {
  terms <- c()

  for (i in 1:2) {
    for (j in 1:2) {
      terms <- c(terms, paste0("I(", var1, "^", i, " * ", var2, "^", j, ")"))
    }
  }

  return(terms)
}

int_X66_X21 <- generate_interactions2("X66", "X21")
int_X69_X21 <- generate_interactions2("X69", "X21")
int_X21_X3 <- generate_interactions2("X21", "X3")

interaction_list <- c(int_X66_X21, int_X69_X21, int_X21_X3)
length(interaction_list) # should be 12

```

```

## [1] 12
# This is lm6, not lm7 to check whether higher interactions are better.

```

```

base_string <- paste(
  "y ~",
  "poly(X66,5,raw=TRUE) +",
  "poly(X69,3,raw=TRUE) +",
  "poly(X21,14,raw=TRUE) +",
  "poly(X3,3,raw=TRUE) +",
  "X19 + X111 + X59 "
)

```

```

test_high_interactions <- function(base_string, interaction_list, train) {

  results <- data.frame()

  for (int in interaction_list) {

    form_string <- paste(base_string, "+", int)
    form <- as.formula(form_string)

    cv <- cv_rmse(form, train)$mean_RMSE

    results <- rbind(
      results,
      data.frame(
        interaction = int,
        RMSE = cv
      )
    )
  }

  return(results)
}

results_deg <- test_high_interactions(base_string, interaction_list, train)

# Sort by RMSE
sorted_interactions <- results_deg[order(results_deg$RMSE), ]

rmse_lm6 <- cv_rmse(lm6, train)$mean_RMSE

better <- sorted_interactions[ sorted_interactions$RMSE < rmse_lm6 , ]

cat("\n==== Interactions BETTER than lm6 ===\n")

## 
## === Interactions BETTER than lm6 ===

print(better)

##           interaction      RMSE
## 4  I(X66^2 * X21^2) 0.4626677
## 3  I(X66^2 * X21^1) 0.4691350
## 2  I(X66^1 * X21^2) 0.4708740
## 1  I(X66^1 * X21^1) 0.4737751
## 6  I(X69^1 * X21^2) 0.4819969
## 5  I(X69^1 * X21^1) 0.4850535
## 10 I(X21^1 * X3^2) 0.4902256
## 12 I(X21^2 * X3^2) 0.4910736
## 9   I(X21^1 * X3^1) 0.4928832

better_terms <- c(
  "I(X66^2 * X21^2)",

```

```

    "I(X66^2 * X21^1)",
    "I(X66^1 * X21^2)",
    "I(X66^1 * X21^1)",
    "I(X69^1 * X21^2)",
    "I(X69^1 * X21^1)",
    "I(X21^1 * X3^2)",
    "I(X21^2 * X3^2)",
    "I(X21^1 * X3^1)"
  )

# best subset 2^9=512 subsets, it's still doable

best_subset <- function(base_string, terms, train) {

  best_rmse <- Inf
  best_model <- NULL
  best_terms <- NULL

  k <- length(terms)

  for (size in 0:k) {

    if (size == 0) {
      subsets <- list(character(0))
    } else {
      subsets <- combn(terms, size, simplify = FALSE)
    }

    for (subset in subsets) {

      # Build formula
      if (length(subset) == 0) {
        form_string <- base_string
      } else {
        form_string <- paste(base_string, "+", paste(subset, collapse = " + "))
      }

      rmse <- cv_rmse(as.formula(form_string), train)$mean_RMSE

      if (rmse < best_rmse) {
        best_rmse <- rmse
        best_model <- form_string
        best_terms <- subset
        message("New best RMSE: ", rmse,
               " | terms: ", paste(subset, collapse=""))
      }
    }
  }

  list(
    best_rmse = best_rmse,
    best_terms = best_terms,
    best_model = best_model
  )
}

```

```

    )
}

result_bs <- best_subset(base_string, better_terms, train)

result_bs$best_rmse

## [1] 0.4513251
result_bs$best_terms

## [1] "I(X66^2 * X21^2)" "I(X69^1 * X21^2)" "I(X21^1 * X3^2)"
result_bs$best_model

## [1] "y ~ poly(X66,5,raw=TRUE) + poly(X69,3,raw=TRUE) + poly(X21,14,raw=TRUE) + poly(X3,3,raw=TRUE) +"
lm8 = lm(y~poly(X66, 5, raw=TRUE) +
          poly(X69, 3, raw=TRUE) +
          poly(X21, 14, raw=TRUE) +
          poly(X3, 3, raw=TRUE) +
          +I(X66^2 * X21^2) + I(X69^1 * X21^2) + I(X21^1 * X3^2) +
          +X19 + X111+X59)
summary(lm8)

##
## Call:
## lm(formula = y ~ poly(X66, 5, raw = TRUE) + poly(X69, 3, raw = TRUE) +
##     poly(X21, 14, raw = TRUE) + poly(X3, 3, raw = TRUE) + +I(X66^2 *
##     X21^2) + I(X69^1 * X21^2) + I(X21^1 * X3^2) + +X19 + X111 +
##     X59)
##
## Residuals:
##       Min      1Q      Median      3Q      Max
## -1.74511 -0.21525 -0.01107  0.21523  1.60252
##
## Coefficients: (1 not defined because of singularities)
##             Estimate Std. Error t value Pr(>|t|)
## (Intercept) 5.927e-01 3.158e+00  0.188 0.851298
## poly(X66, 5, raw = TRUE)1 -1.897e-02 8.432e-02 -0.225 0.822163
## poly(X66, 5, raw = TRUE)2  4.233e-04 1.267e-03  0.334 0.738643
## poly(X66, 5, raw = TRUE)3  4.981e-05 5.061e-05  0.984 0.325937
## poly(X66, 5, raw = TRUE)4  5.116e-06 3.505e-06  1.460 0.145521
## poly(X66, 5, raw = TRUE)5 -1.877e-07 6.859e-08 -2.737 0.006608 **
## poly(X69, 3, raw = TRUE)1  1.445e-02 8.379e-02  0.172 0.863243
## poly(X69, 3, raw = TRUE)2 -1.461e-03 4.890e-04 -2.989 0.003060 **
## poly(X69, 3, raw = TRUE)3 -2.572e-05 1.612e-05 -1.596 0.111764
## poly(X21, 14, raw = TRUE)1 2.911e+00 1.563e+00  1.862 0.063708 .
## poly(X21, 14, raw = TRUE)2 -7.117e+00 3.115e+00 -2.284 0.023124 *
## poly(X21, 14, raw = TRUE)3  6.487e+00 2.518e+00  2.576 0.010522 *
## poly(X21, 14, raw = TRUE)4 -3.137e+00 1.102e+00 -2.848 0.004742 **
## poly(X21, 14, raw = TRUE)5  9.107e-01 2.954e-01  3.082 0.002266 **
## poly(X21, 14, raw = TRUE)6 -1.696e-01 5.189e-02 -3.268 0.001222 **
## poly(X21, 14, raw = TRUE)7  2.107e-02 6.182e-03  3.407 0.000756 ***
## poly(X21, 14, raw = TRUE)8 -1.777e-03 5.067e-04 -3.507 0.000531 ***

```

```

## poly(X21, 14, raw = TRUE)9  1.016e-04  2.842e-05  3.576 0.000414 ***
## poly(X21, 14, raw = TRUE)10 -3.833e-06  1.058e-06 -3.621 0.000350 ***
## poly(X21, 14, raw = TRUE)11  8.848e-08  2.424e-08  3.650 0.000315 ***
## poly(X21, 14, raw = TRUE)12 -1.004e-09  2.740e-10 -3.665 0.000297 ***
## poly(X21, 14, raw = TRUE)13      NA        NA        NA        NA
## poly(X21, 14, raw = TRUE)14  8.074e-14  2.199e-14  3.672 0.000290 ***
## poly(X3, 3, raw = TRUE)1    -1.144e-01  8.234e-02 -1.389 0.165911
## poly(X3, 3, raw = TRUE)2    2.077e-03  9.398e-04  2.210 0.027941 *
## poly(X3, 3, raw = TRUE)3    -1.155e-05  4.374e-06 -2.641 0.008752 **
## I(X66^2 * X21^2)           1.500e-06  1.248e-06  1.202 0.230290
## I(X69^1 * X21^2)           8.286e-05  3.476e-05  2.384 0.017816 *
## I(X21^1 * X3^2)            -1.178e-05  2.706e-06 -4.353 1.91e-05 ***
## X19                          8.307e-02  2.346e-02  3.541 0.000469 ***
## X111                         1.012e+00  1.688e-01  5.997 6.44e-09 ***
## X59                          -5.360e-02  3.746e-02 -1.431 0.153591
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.4324 on 269 degrees of freedom
## Multiple R-squared:  0.8037, Adjusted R-squared:  0.7818
## F-statistic: 36.72 on 30 and 269 DF,  p-value: < 2.2e-16
cv_rmse(lm8,train)

```

```

## $mean_RMSE
## [1] 0.4513251
##
## $sd_RMSE
## [1] 0.0767361
##
## $all_RMSE
## [1] 0.4425578 0.5221690 0.5043998 0.4000440 0.4901194 0.3625572 0.3188291
## [8] 0.5083940 0.4101056 0.5540755

```

Now our model is quite complex. Let's see if removing any of the term could reduce cv rmse.

```

blocks <- c(
  "poly(X66,5,raw=TRUE)",
  "poly(X69,3,raw=TRUE)",
  "poly(X21,14,raw=TRUE)",
  "poly(X3,3,raw=TRUE)",
  "I(X66^2 * X21^2)",
  "I(X69^1 * X21^2)",
  "I(X21^1 * X3^2)",
  "X19",
  "X111",
  "X59"
)

best_subset_blocks <- function(block_terms, train) {

  best_rmse  <- Inf
  best_model <- NULL
  best_terms <- NULL

  k <- length(block_terms)

```

```

for (size in 1:k) {# or 0:k if you want intercept-only too
  subsets <- combn(block_terms, size, simplify = FALSE)

  for (subset in subsets) {
    rhs <- paste(subset, collapse = " + ")
    form_string <- paste("y ~", rhs)

    rmse <- cv_rmse(as.formula(form_string), train)$mean_RMSE

    if (rmse < best_rmse) {
      best_rmse <- rmse
      best_model <- form_string
      best_terms <- subset
      message("New best RMSE: ", rmse,
              " | terms: ", paste(subset, collapse = " , "))
    }
  }
}

list(
  best_rmse = best_rmse,
  best_terms = best_terms,
  best_model = best_model
)
}

result_blocks <- best_subset_blocks(blocks, train)
result_blocks$best_rmse

## [1] 0.4496267
result_blocks$best_terms

## [1] "poly(X66,5,raw=TRUE)"  "poly(X69,3,raw=TRUE)"  "poly(X21,14,raw=TRUE)"
## [4] "poly(X3,3,raw=TRUE)"    "I(X66^2 * X21^2)"    "I(X69^1 * X21^2)"
## [7] "I(X21^1 * X3^2)"       "X19"                  "X111"
result_blocks$best_model

## [1] "y ~ poly(X66,5,raw=TRUE) + poly(X69,3,raw=TRUE) + poly(X21,14,raw=TRUE) + poly(X3,3,raw=TRUE) +
lm9 = lm(y ~ poly(X66,5,raw=TRUE) +
          poly(X69,3,raw=TRUE) +
          poly(X21,14,raw=TRUE) +
          poly(X3,3,raw=TRUE) +
          I(X66^2 * X21^2) +
          I(X69^1 * X21^2) +
          I(X21^1 * X3^2) +
          X19 + X111)
summary(lm9)

##
## Call:
## lm(formula = y ~ poly(X66, 5, raw = TRUE) + poly(X69, 3, raw = TRUE) +
##     poly(X21, 14, raw = TRUE) + poly(X3, 3, raw = TRUE) + I(X66^2 *
##     X21^2) + I(X69^1 * X21^2) + I(X21^1 * X3^2) + X19 + X111)

```

```

##
## Residuals:
##      Min     1Q Median     3Q    Max
## -1.74289 -0.21825 -0.00186  0.21783  1.59882
##
## Coefficients: (1 not defined because of singularities)
##                               Estimate Std. Error t value Pr(>|t|)
## (Intercept)                 5.605e-01  3.165e+00   0.177 0.859558
## poly(X66, 5, raw = TRUE)1 -1.892e-02  8.448e-02  -0.224 0.823001
## poly(X66, 5, raw = TRUE)2  5.020e-04  1.269e-03   0.396 0.692613
## poly(X66, 5, raw = TRUE)3  5.051e-05  5.070e-05   0.996 0.320090
## poly(X66, 5, raw = TRUE)4  4.795e-06  3.504e-06   1.368 0.172375
## poly(X66, 5, raw = TRUE)5 -1.813e-07  6.857e-08  -2.643 0.008686 **
## poly(X69, 3, raw = TRUE)1  1.292e-02  8.394e-02   0.154 0.877749
## poly(X69, 3, raw = TRUE)2 -1.418e-03  4.889e-04  -2.899 0.004048 **
## poly(X69, 3, raw = TRUE)3 -2.432e-05  1.612e-05  -1.509 0.132549
## poly(X21, 14, raw = TRUE)1 2.842e+00  1.566e+00   1.815 0.070585 .
## poly(X21, 14, raw = TRUE)2 -6.912e+00  3.118e+00  -2.217 0.027474 *
## poly(X21, 14, raw = TRUE)3  6.286e+00  2.519e+00   2.495 0.013180 *
## poly(X21, 14, raw = TRUE)4 -3.035e+00  1.101e+00  -2.756 0.006252 **
## poly(X21, 14, raw = TRUE)5  8.801e-01  2.952e-01   2.981 0.003135 **
## poly(X21, 14, raw = TRUE)6 -1.637e-01  5.183e-02  -3.159 0.001765 **
## poly(X21, 14, raw = TRUE)7  2.030e-02  6.171e-03   3.290 0.001134 **
## poly(X21, 14, raw = TRUE)8 -1.710e-03  5.055e-04  -3.383 0.000823 ***
## poly(X21, 14, raw = TRUE)9  9.765e-05  2.834e-05   3.446 0.000660 ***
## poly(X21, 14, raw = TRUE)10 -3.677e-06  1.055e-06  -3.486 0.000572 ***
## poly(X21, 14, raw = TRUE)11  8.476e-08  2.415e-08   3.510 0.000526 ***
## poly(X21, 14, raw = TRUE)12 -9.606e-10  2.728e-10  -3.521 0.000505 ***
## poly(X21, 14, raw = TRUE)13        NA       NA       NA       NA
## poly(X21, 14, raw = TRUE)14  7.699e-14  2.187e-14   3.520 0.000506 ***
## poly(X3, 3, raw = TRUE)1  -1.176e-01  8.247e-02  -1.426 0.155079
## poly(X3, 3, raw = TRUE)2  2.172e-03  9.393e-04   2.312 0.021507 *
## poly(X3, 3, raw = TRUE)3 -1.219e-05  4.360e-06  -2.795 0.005559 **
## I(X66^2 * X21^2)           1.381e-06  1.247e-06   1.107 0.269180
## I(X69^1 * X21^2)           9.092e-05  3.436e-05   2.646 0.008626 **
## I(X21^1 * X3^2)            -1.297e-05  2.579e-06  -5.029 8.99e-07 ***
## X19                         8.607e-02  2.341e-02   3.677 0.000284 ***
## X111                        1.016e+00  1.691e-01   6.008 6.07e-09 ***
##
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.4332 on 270 degrees of freedom
## Multiple R-squared:  0.8022, Adjusted R-squared:  0.781
## F-statistic: 37.77 on 29 and 270 DF,  p-value: < 2.2e-16
cv_rmse(lm9,train)

## $mean_RMSE
## [1] 0.4496267
##
## $sd_RMSE
## [1] 0.07273673
##
## $all_RMSE
## [1] 0.4415505 0.4814945 0.5105339 0.4084901 0.4914399 0.3642895 0.3333033

```

```

## [8] 0.5133146 0.3948952 0.5569552
# Now we have finished everything with lm. Let's try it on Kaggle.

test <- read.csv("test_predictors.csv")

pred <- predict(lm9, newdata = test)
sample <- read.csv("SampleSubmission.csv")

submission <- data.frame(
  id = sample$id,
  y   = pred
)

write.csv(submission, "LM9_Submission.csv", row.names=FALSE)

cat("Saved: LM9_Submission.csv\n")

## Saved: LM9_Submission.csv

Score on Kaggle: Private Score: 0.52147 Public Score: 0.52681

Now let's use GAM to get lower RMSE

s: smoothing splines ti: smooth interactions

library(mgcv)

# This is similar with LM9,
# but we change the one variable polynomials to smoothing splines
gam1_formula <- y ~
  s(X66, k = 12) +
  s(X69, k = 12) +
  s(X21, k = 18) +
  s(X3, k = 10) +
  I(X66^2 * X21^2) +
  I(X69^1 * X21^2) +
  I(X21^1 * X3^2) +
  X19 + X111

# Fit GAM1
gam1 <- gam(gam1_formula, data=train, method="REML", select=TRUE)

summary(gam1)

##
## Family: gaussian
## Link function: identity
##
## Formula:
## y ~ s(X66, k = 12) + s(X69, k = 12) + s(X21, k = 18) + s(X3,
##       k = 10) + I(X66^2 * X21^2) + I(X69^1 * X21^2) + I(X21^1 *
##       X3^2) + X19 + X111
##
## Parametric coefficients:
##                         Estimate Std. Error t value Pr(>|t|)
## (Intercept)      -7.287e-01  2.235e-01 -3.261 0.001254 **
```

```

## I(X66^2 * X21^2)  1.124e-06  1.156e-06  0.972 0.331684
## I(X69^1 * X21^2)  9.951e-05  3.153e-05  3.156 0.001783 **
## I(X21^1 * X3^2)  -1.369e-05  2.428e-06  -5.636 4.38e-08 ***
## X19                 8.606e-02  2.313e-02  3.720 0.000242 ***
## X111                9.929e-01  1.649e-01  6.020 5.70e-09 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Approximate significance of smooth terms:
##          edf Ref.df      F  p-value
## s(X66)    6.683     11  5.353 < 2e-16 ***
## s(X69)    3.583     11  2.948 < 2e-16 ***
## s(X21)   11.606     17 12.443 < 2e-16 ***
## s(X3)     3.049      9  1.694 0.000354 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## R-sq.(adj) =  0.795  Deviance explained = 81.5%
## -REML = 247.57  Scale est. = 0.17576 n = 300
# Create cv_gam function
cv_gam <- function(formula, data, K = 10, seed = 1) {
  set.seed(seed)

  n <- nrow(data)
  folds <- sample(rep(1:K, length.out = n))

  rmse <- numeric(K)

  for (k in 1:K) {
    train_fold <- data[folds != k, ]
    test_fold <- data[folds == k, ]

    # Fit GAM
    fit <- gam(formula, data = train_fold, method = "REML")

    # Predict
    pred <- predict(fit, newdata = test_fold)

    # RMSE
    rmse[k] <- sqrt(mean((pred - test_fold$y)^2))
  }

  return(list(
    mean_RMSE = mean(rmse),
    sd_RMSE   = sd(rmse),
    all_RMSE  = rmse
  ))
}

cv_gam1 = cv_gam(gam1_formula,train)
cv_gam1

## $mean_RMSE
## [1] 0.4449057
##

```

```

## $sd_RMSE
## [1] 0.07555486
##
## $all_RMSE
## [1] 0.4209355 0.4885687 0.5117365 0.4091546 0.5194580 0.3587830 0.3304465
## [8] 0.4874085 0.3754565 0.5471092

```

Then, we change the interactions to smooth interactions.

```

gam2_formula <- y ~
  s(X66, k = 12) +
  s(X69, k = 12) +
  s(X21, k = 18) +
  s(X3, k = 10) +
  ti(X66, X21, k=8) +
  ti(X69, X21, k=8) +
  ti(X21, X3, k=8) +
  X19 + X111

gam2 <- gam(gam2_formula, data=train, method="REML", select=TRUE)

summary(gam2)

##
## Family: gaussian
## Link function: identity
##
## Formula:
## y ~ s(X66, k = 12) + s(X69, k = 12) + s(X21, k = 18) + s(X3,
##       k = 10) + ti(X66, X21, k = 8) + ti(X69, X21, k = 8) + ti(X21,
##       X3, k = 8) + X19 + X111
##
## Parametric coefficients:
##             Estimate Std. Error t value Pr(>|t|)
## (Intercept) -1.23535   0.16809 -7.349 2.62e-12 ***
## X19          0.08444   0.02219  3.805 0.000177 ***
## X111         1.01308   0.15593  6.497 4.19e-10 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Approximate significance of smooth terms:
##             edf Ref.df      F p-value
## s(X66)      6.3779    11  5.901 <2e-16 ***
## s(X69)      3.6854    11  2.475 <2e-16 ***
## s(X21)     11.7302    17 16.388 <2e-16 ***
## s(X3)        3.2428     9  3.696 <2e-16 ***
## ti(X66,X21) 5.6233    49  1.504 <2e-16 ***
## ti(X69,X21) 0.9208    49  0.022   0.216
## ti(X21,X3)  7.1478    49  0.847 <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## R-sq.(adj) =  0.818  Deviance explained = 84.3%
## -REML = 213.17  Scale est. = 0.15568 n = 300

```

```

cv_gam2 = cv_gam(gam2_formula,train)
cv_gam2

## $mean_RMSE
## [1] 0.4325536
##
## $sd_RMSE
## [1] 0.07460207
##
## $all_RMSE
## [1] 0.4027576 0.4780151 0.4697028 0.4321348 0.4904316 0.3457749 0.2967451
## [8] 0.5149694 0.3797129 0.5152916
cv_gam2$mean_RMSE <- cv_gam1$mean_RMSE

## [1] TRUE

```

From the summary function, the $ti(X69, X21)$ p-value is quite suspicious. Let's try removing it and check the RMSE.

```

gam3_formula <- y ~
  s(X66, k = 12) +
  s(X69, k = 12) +
  s(X21, k = 18) +
  s(X3, k = 10) +
  ti(X66, X21, k=8) +
  ti(X21, X3, k=8) +
  X111 + X19

gam3 <- gam(gam3_formula, data=train, method="REML", select=TRUE)

summary(gam3)

```

```

##
## Family: gaussian
## Link function: identity
##
## Formula:
## y ~ s(X66, k = 12) + s(X69, k = 12) + s(X21, k = 18) + s(X3,
##       k = 10) + ti(X66, X21, k = 8) + ti(X21, X3, k = 8) + X111 +
##       X19
##
## Parametric coefficients:
##             Estimate Std. Error t value Pr(>|t|)
## (Intercept) -1.23700   0.16826 -7.352 2.57e-12 ***
## X111         1.01484   0.15607  6.502 4.06e-10 ***
## X19          0.08448   0.02222  3.802 0.000179 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Approximate significance of smooth terms:
##             edf Ref.df      F p-value
## s(X66)      6.396    11  6.064 <2e-16 ***
## s(X69)      3.692    11  2.534 <2e-16 ***
## s(X21)     11.725    17 16.317 <2e-16 ***
## s(X3)       3.225     9  3.640 <2e-16 ***

```

```

## ti(X66,X21) 6.085      49  1.839  <2e-16 ***
## ti(X21,X3)   7.297      49  0.860  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## R-sq.(adj) =  0.818    Deviance explained = 84.3%
## -REML = 213.2  Scale est. = 0.15601 n = 300
cv_gam3 = cv_gam(gam3_formula,train)
cv_gam3

## $mean_RMSE
## [1] 0.4308959
##
## $sd_RMSE
## [1] 0.07402865
##
## $all_RMSE
## [1] 0.4033649 0.4767075 0.4680794 0.4332528 0.4886535 0.3405254 0.2961775
## [8] 0.5076227 0.3802012 0.5143739
cv_gam3$mean_RMSE < cv_gam2$mean_RMSE

```

[1] TRUE

X19 is a discrete variable, treating it as a categorical variable may improve the GAM's performance.

```

train$X19_f <- factor(train$X19)
test$X19_f  <- factor(test$X19, levels = levels(train$X19_f))

gam4_formula <- y ~
  s(X66, k = 12) +
  s(X69, k = 12) +
  s(X21, k = 18) +
  s(X3, k = 10) +
  ti(X66, X21, k=8) +
  ti(X21, X3, k=8) +
  X111 + X19_f

gam4 <- gam(gam4_formula, data=train, method="REML", select=TRUE)

summary(gam4)

##
## Family: gaussian
## Link function: identity
##
## Formula:
## y ~ s(X66, k = 12) + s(X69, k = 12) + s(X21, k = 18) + s(X3,
##       k = 10) + ti(X66, X21, k = 8) + ti(X21, X3, k = 8) + X111 +
##       X19_f
##
## Parametric coefficients:
##             Estimate Std. Error t value Pr(>|t|)
## (Intercept) -1.25771   0.17926 -7.016 2.03e-11 ***

```

```

## X111      1.00641   0.15544   6.474 4.82e-10 ***
## X19_f1    0.03075   0.12121   0.254  0.79993
## X19_f3    0.43893   0.15439   2.843  0.00483 **
## X19_f4    0.29296   0.12571   2.330  0.02057 *
##
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Approximate significance of smooth terms:
##          edf Ref.df   F p-value
## s(X66)     6.321     11  5.790 <2e-16 ***
## s(X69)     3.684     11  2.711 <2e-16 ***
## s(X21)    11.709     17 16.374 <2e-16 ***
## s(X3)      3.391      9  3.708 <2e-16 ***
## ti(X66,X21) 5.869     49  1.874 <2e-16 ***
## ti(X21,X3)  7.819     49  0.931 <2e-16 ***
##
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## R-sq.(adj) =  0.821  Deviance explained = 84.7%
## -REML = 211.63  Scale est. = 0.15312 n = 300
cv_gam4 = cv_gam(gam4_formula,train)
cv_gam4

## $mean_RMSE
## [1] 0.4270584
##
## $sd_RMSE
## [1] 0.07746553
##
## $all_RMSE
## [1] 0.4060467 0.4646716 0.4546878 0.4237394 0.5037809 0.3228253 0.2858733
## [8] 0.4970759 0.3922138 0.5196695
cv_gam4$mean_RMSE <- cv_gam3$mean_RMSE

## [1] TRUE
pred <- predict(gam4, newdata = test)
sample <- read.csv("SampleSubmission.csv")

submission <- data.frame(
  id = sample$id,
  y   = pred
)

write.csv(submission, "GAM4_Submission.csv", row.names=FALSE)

cat("Saved: GAM4_Submission.csv\n")

## Saved: GAM4_Submission.csv

Score on Kaggle: Private Score: 0.48633 Public Score: 0.49299

```